Appendix B. Example Data Sheet - Commercial Sensor

OVERVIEW	35
Compliance and Certifications	35
Description	35
Features	35
Use Cases	35
MODEL CHARACTERISTICS	36
Software Flow Diagram	37
Dataset Nutrition Label	38
IoT Security and Privacy Label	39
Machine Learning Model Specification	41
Person Detection Model	41
Face Identification Model	41
Performance Analysis	43
Environmental Sensitivity	44
Demographic biases	44
HARDWARE	45
Hardware Details	46
I2C Protocol	47
Device Diagrams	48
Bill of Materials	50
Environmental Impact	51
Acronyms	52
Glossary	53

OVERVIEW

Person Sensor V1.0 SEN-21231



Compliance and Certifications

The person detection sensor complies with essential industry standards and regulations, including RoHS for environmental safety and GDPR for protecting individual privacy. The sensor has been audited by Kodelski Security for security and privacy implications.

Description

The Person Sensor from Useful Sensors is a small, low-cost hardware module that detects nearby peoples' faces, and returns information about how many there are, where they are relative to the device, and performs facial recognition. It is designed to be used as an input to a larger system, for example to wake up a kiosk display from sleep mode when somebody approaches, mute a microphone when nobody is present, or orient a fan so it's always pointing at the nearest person.

The sensor has a small form factor and utilizes a monochrome camera with a field of view of 640 x 480 (VGA). The input voltage for the sensor is 3.3V and the typical operating current for the sensor is 40 mA. The sensor communicates via I2C/Qwiic mode, conforming to SparkFun Qwiic electrical/mechanical specifications, and has a maximum cable length of 1 m at 400 kb/s. Longer cables can be used at lower data rates. The sensor has a maximum data rate of 400 kb/s.

Features

- Real-time person + head pose tracking with on-device ML
- Real-time person identification with on-devicce ML
- Low power consumption
- Onboard camera
- Small form factor: 22 x 20 x 10 mm
- I2C serial communication
- Lead-free

Use Cases

- Security
- Home automation
- Consumer appliances

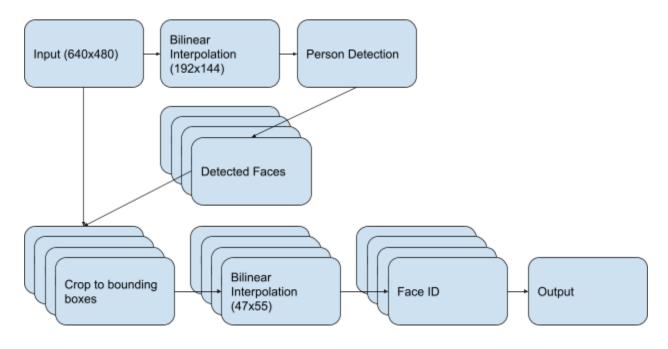
MODEL CHARACTERISTICS

Software Flow Diagram

8-bit grayscale images (640x480) are resized to 192x144 and passed into a RetinaFace model trained to detect human faces. This model outputs a list of faces with coordinates for a bounding box around each face as well as five key facial landmarks. If identity is enabled, bounded faces are cropped out of the original image and rescaled to 47x55 and passed into a DeepID model to generate an embedding. This embedding is compared with saved facial IDs, and the nearest ID is returned along with the bounding box and information about whether the person is facing the sensor. Output information is communicated via Qwiic interface to the application processor.

	Person Detection Model	Person ID Model	
Architecture	RetinaFace	DeepID	
Framework	TFLite Micro	TFLite Micro	
Validation Set	20%	-	
Quantization	int8	int8	
Inference Time	140 ms	125 ms	
Peak RAM Usage	442.6 kB	189 kB	
Flash Usage	449 kB	397 kB	

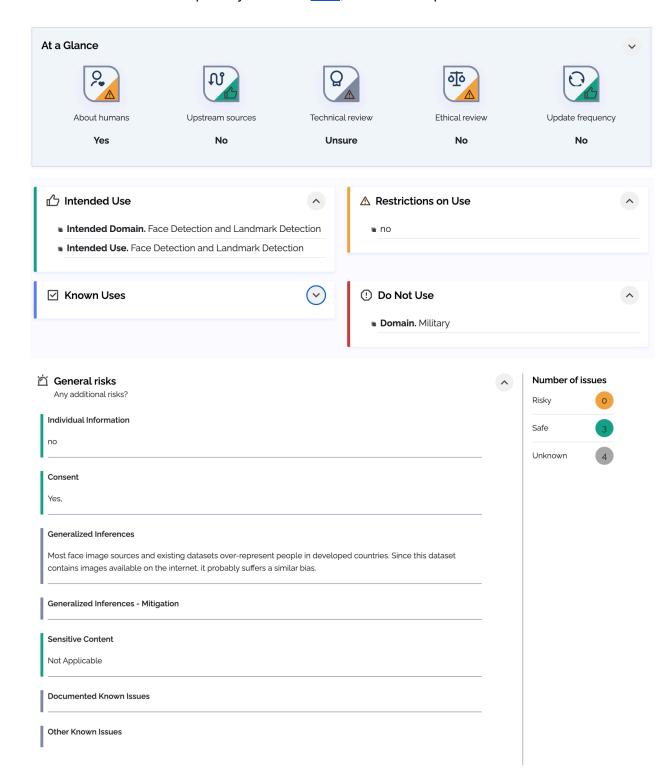
Person Sensor Software Flow



Jan 2024 - Rev. 2 37 Person Sensor V1.0

Dataset Nutrition Label

The data nutrition label is publicly available here, with some important features outlined below.



Security & Privacy Overview Useful Sensors Person Sensor V1.0 Firmware version: Not disclosed - updated on: 2023-05-03 The device was manufactured in: China Security updates No security updates (i) Not disclosed Access control Security (i) **Mechanisms** Sensor data collection Visual **Physiological** Camera Sensor type Providing and improving **Purpose** device functions Data stored on No device storage the device **Data Practices** Data stored in No cloud storage the cloud Data shared with Not shared Data sold to Not sold Other collected Privacy policy Not disclosed (i) **Detailed Security & Privacy Label:** Not disclosed **More Information** CMU IoT Security and Privacy Label CISPL 1.0 iotsecurityprivacy.org

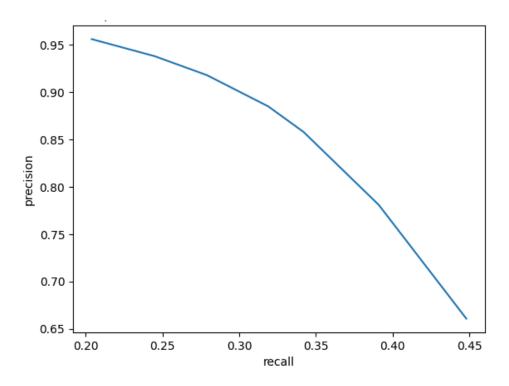
Security & Privacy Details					
Person Sensor V1.0	ot disclosed - updated or				
	Security updates		(0	No security updates Sensor is a standatone unit. For security and privacy, the firmware cannot be changed and only sensor outputs
	Access control			(i)	are available. Not disclosed
	Security oversight			0	Audits performed by third-party security auditors
Д	Ports and protocols			(i)	Third party security audit performed by Kudelski Security Not disclosed
Security Mechanisms	Hardware safety			<u> </u>	Not disclosed
,,,,,,,,,,,,,,,,,,,,,,,,,,,,,,,,,,,,,,,	Software safety			ı	Not disclosed
	Personal safety			0	Not disclosed
	Vulnerability disclosure a	and management		0	Not disclosed
	Software and hardware o	omposition list		0	Not disclosed
	Encryption and key mana	gement		0	Not disclosed
© Data Practices	Sensor data collection Sensor type Data collection frequency Purpose Data stored on the device Local data retention time Data stored in the cloud Cloud data retention time Data shared with Data sharing frequency Data sold to Other collected data	Visual Camera Continuous Providing and improving device functions No device storage No retention No cloud storage No retention Not shared Not shared Not sold None			
	Data linkage		((i)	Not disclosed
	What will be Inferred fr		(0	Not disclosed
	Special data handling p	practices for children	(0	Not disclosed
	In Compliance with		(0	Not disclosed
	Privacy policy		((i)	Not disclosed
	Call Useful Sensors with	your questions at	(①	1 805 813 7571
	Email Useful Sensors with your questions at		(0	contact@usefulsensors.com
8	Functionality when offline	6	((i)	Full functionality on offline mode
More Information	Functionality with no data	a processing	((i)	Not applicable
	Physical actuations and t	riggers		1	Not disclosed
	Compatible platforms		(1	Not disclosed
CMU IoT Security and	Privacy Label CISPL 1.0 iots	ecurityprivacy.org			

Jan 2024 - Rev. 2 40 Person Sensor V1.0

Machine Learning Model Specification

Person Detection Model

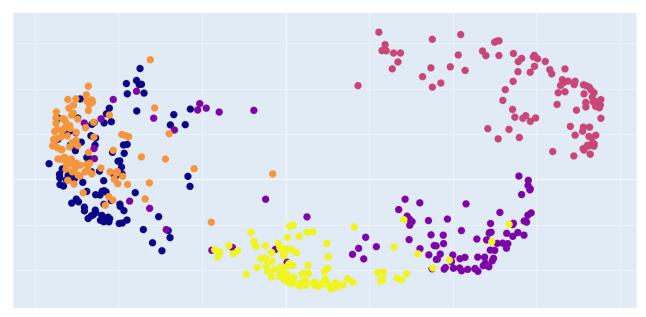
The person detection model was trained on a proprietary dataset of ~30,000 images with 300k labeled faces and five facial landmarks per face. The model input is a 192x144 raw image in 8-bit grayscale format, equivalent to 27,648 features. The training process was performed until the model accuracy ceased to improve. Final model performance achieved a precision of 91.8% on the test set, using a threshold of 0.7. The precision-recall curve of the model on the test set is shown below. The model was quantized to 8-bit integer using post training quantization using the Tensorflow Lite converter and is deployed using the Tensorflow Lite Micro runtime.



Face Identification Model

The face identification model was fine-tuned on a proprietary dataset encompassing ~4000 images across 5 identities captured using the sensor camera module. The input to the model is a 47x55 raw image in 8-big grayscale format, equivalent to 2,585 features. To produce the best separation between faces a dense classification layer was added to the model, and several iterations of freezing either the classification layer or the model was used to achieve a higher accuracy on the fine-tuning dataset. Finally, the classification layer was removed and embedding separation was evaluated using Principal Component Analysis (PCA) in three dimensions.

Jan 2024 - Rev. 2 41 Person Sensor V1.0



Each color represents one of five unique identities in the validation dataset. Distances between points indicate approximate distances between embeddings simplified to 2-D space.

The model was quantized to 8-bit integer using post training quantization using the Tensorflow Lite converter and is deployed using the Tensorflow Lite Micro runtime. On the sensor, the embedding generated by the Face ID model is compared against registered faces, and if a face with similar enough features is found, that identity is used. Otherwise, an identity of -1 is returned to indicate that no registered identity was found.

Jan 2024 - Rev. 2 42 Person Sensor V1.0

Performance Analysis

The end-to-end performance of the person detection sensor model was tested through an experimental study conducted in the Science and Engineering Complex (SEC) at Harvard University. The study involved 40 participants and evaluated the accuracy of the model under different lighting conditions using three identical sensors.

The study room measured 25 x 31 x 10 ft and contained 32 ceiling lights that were uniformly distributed in a 4 x 8 grid. The lighting conditions were captured quantitatively for each participant using a <u>Lux LCD</u> <u>Illuminance Meter</u> (Precision Vision, Inc.) and a <u>C-800-U Spectrometer</u> (Sekonic Corporation).

The sensors were mounted on a wooden board affixed to the wall at a height of 1.5 m above the ground. The participants were evaluated at three different distances (1.5 m, 4.5 m, and 7.5 m) from the sensors under each lighting condition. The ambient lighting in the room was provided by artificial lights, and blackout curtains were used to block the ambient lighting from outside.

The lighting levels were controlled using a dimmer switch that had three levels of operation, corresponding to 208±31, 584±51, and 1149±59 lux, respectively. When the lights were turned off, the illuminance meter gave a reading of zero lux. When all the lights were turned on at full strength, the sensor gave an average reading of 1149 lux. The color temperature of the lighting was measured to be 5600 K, corresponding to white light. Colored tape was placed on the ground to demarcate the locations where participants should stand during the experiment (i.e., 1.5, 4.5, and 7.5 m from the sensor array).

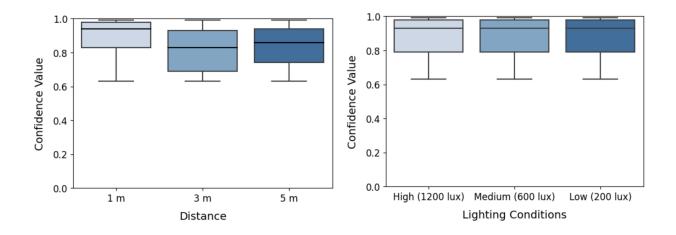
Before entering the study environment, the participants were asked to provide their gender identity and evaluate their skin tone according to the Monk Skin Tone (MST) Scale to evaluate algorithmic bias. The study evaluated algorithmic bias by bucketing skin tone into three categories: light (MST 0-4), medium (MST 5-7), and dark (MST 8-10). At each location and lighting condition, ten readings were taken from each sensor and averaged.

Participants were recruited using flyers, and those interested filled out a Study Interest Form. Upon arrival, participants signed a Consent Form indicating their willingness to participate in the study. The accuracy of the model is provided in the following graphs as a function of lighting condition, distance, gender identity, and skin tone. Overall, 63.2% of the participants were male, and 36.8% were female; the percentage of participants corresponding to each skin tone group was: 47.4% light, 39.4% medium, and 13.2% dark.

Jan 2024 - Rev. 2 43 Person Sensor V1.0

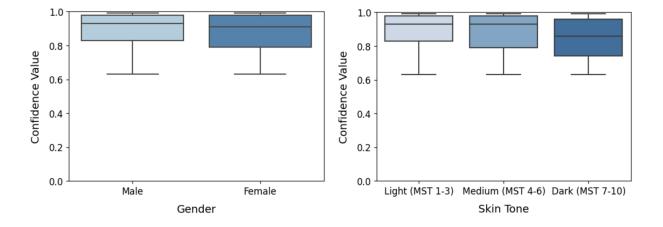
Environmental Sensitivity

The device showed no decrease in performance under decreased lighting conditions. A moderate drop off in performance of around 10% is observed at distances 3-5 meters from the sensor.



Demographic biases

A small gender bias is observed in model performance. A moderate skin tone bias was observed, showing approximately a 10% decrease in the confidence value for individuals with a darker skin tone.



Jan 2024 - Rev. 2 44 Person Sensor V1.0

HARDWARE

Hardware Details

Camera Specifications (see <u>here</u>)			
Field of view (horizontal)	110°		
Color Filter Array	Bayer, Monochrome		
Frame Rate	60FPS @ 48MHz		
Pixel Array (Active/ Effective)	644 x 484 / 640x480		
Electrical Specifications			
Operating Voltage Range (regulator enabled)	3.1V to 3.5V		
Operating Current	40 mA		
Operating Temperature	-20 °C to 85 °C		
Communication Specifications			
I2C/Qwiic mode	Conforms with SparkFun Qwiic electrical/mechanical specifications. https://www.sparkfun.com/qwiic		
Max cable length	1 m		
Max data rate	100 kb/s		
Module Orientation	Text on sensor is upright, up arrow points upwards		
GPIO mode	INT pin is high when person is detected		
Diagnostic LED	Default behavior of green LED on board: illuminates when person detected.		
Data Transfer and Format	See I2C Protocol Table		
I2C Address	0x63		

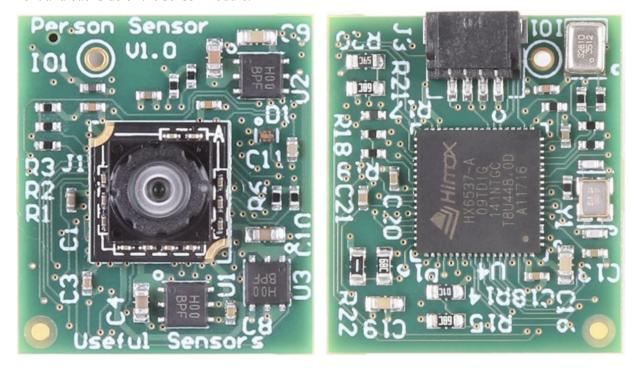
I2C Protocol

Address	Name	Default	Description
0x01	Mode	0x01 (continuous)	Mode. See mode table below.
0x02	Enable ID	0x00 (False)	Enable / Disable the ID model. With this flag set to False, only capture bounding boxes.
0x03	Single shot	0x00	Trigger a single-shot inference. Only works if the sensor is in standby mode.
0x04	Label next ID	0x00	Calibrate the next identified frame as person N, from 0 to 7. If two frames pass with no person, this label is discarded.
0x05	Persist IDs	0x01 (True)	Store any recognized IDs even when unpowered.
0x06	Erase IDs	0x0	Wipe any recognized IDs from storage.
0x07	Debug Mode	0x01 (True)	Whether to enable the LED indicator on the sensor.

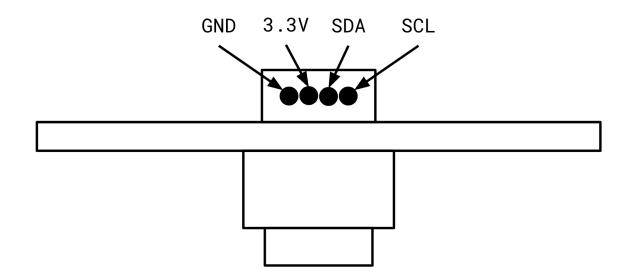
Mode	Name	Description
0x00	Standby	Lowest power mode, sensor is in standby and not capturing.
0x01	Continuous	Capture continuously, setting the GPIO trigger pin to high if a face is detected.

Device Diagrams

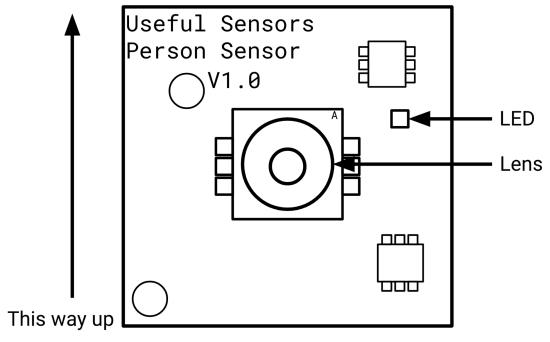
Front and backside of the sensor module.



Qwiic connector interface.

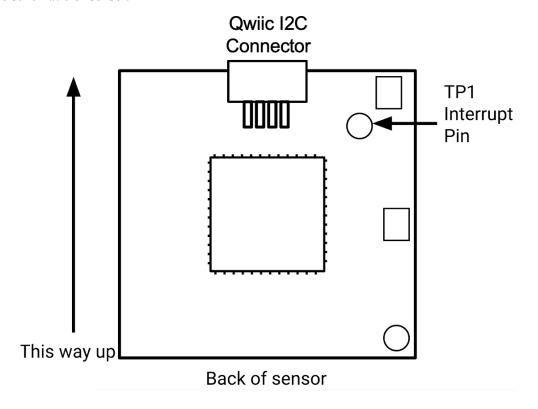


Frontside schematic of sensor.



Front of sensor

Backside schematic of sensor.



Bill of Materials

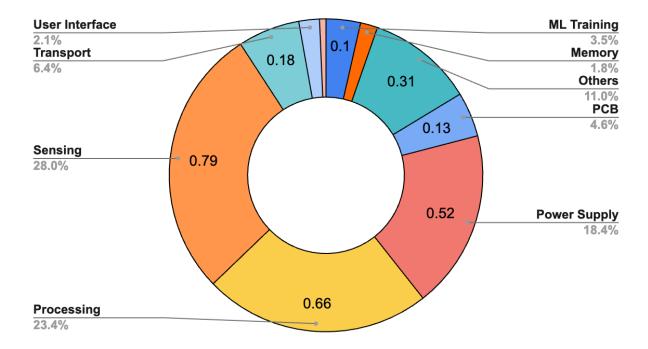
The following is a comprehensive list of materials required to assemble the Person Sensor V1.0 by Useful Sensors. All unit cost values quoted in minimum order quantity of one.

Category In TinyML Calculator	Component	Unit Cost (\$)	Quantity	Manufacturer	Link to Datasheet (if available)
Functional Component	ts	-	<u>-</u>		
V	Himax MCU HX6537/39/40-A	14.50	1	HiMax	https://cdn.sparkfun.com/assets/6/6/7/e/ 8/HX6537-A DS public v01 1 .pdf
V	Camera Module HM0360-MWA	8.81	1	HiMax	https://cdn.sparkfun.com/assets/d/2/9/9/ 7/Pre-HM0360_DS_preliminary_v04_ Ltd1.pdf
V	MEMS Microphone SPH0641LM4H-1	1.05	1	Knowles	https://media.digikey.com/pdf/Data%20 Sheets/Knowles%20Acoustics%20PDF s/SPH0641LM4H-1.pdf
V	Crystal Oscillator ECS-240-10-36-CKM-TR	0.57	1	ECS Inc.	https://ecsxtal.com/store/pdf/ECX-2236 .pdf
Power Circuitry					
	Adjustable Linear Voltage Regulator R1173D001B-TR-FE	1.33	3	Nisshinbo Micro Devices Inc.	https://www.nisshinbo-microdevices.co .jp/en/pdf/datasheet/r1173-ea.pdf
Indication					
V	RGB LED	0.1	1	Harvatek Corporation	https://media.digikey.com/pdf/Data%20 Sheets/Harvatek%20PDFs/B39D3RGB _F6C0001HOU1930.pdf
Connectors					
	Board to Camera OK-10F030-04	1.22	1	AliExpress	N/A
	Qwiic JST SH 4-pin Right Angle Connector	0.40	1	Adafruit	N/A
Passive Components					
✓	Misc resistors	0.01	15	-	N/A
V	Misc capacitors	0.01	17	-	N/A
✓	Mise inductors	0.01	1	-	N/A
	Total	30.97			

Environmental Impact

With the widespread deployment of smart sensors, it is essential to consider and be conscious of the environmental impact such ubiquitous computing may have. Thus another component we advocate to be included in the datasheet is an "environmental impact" section that outlines the device footprint. Using the methodology of [9], we generated a sample of what this section might look like as part of the datasheet for our sensor specifically. We capture the carbon footprint (CO2-eq.) of our ML sensor in the chart below. Due to the limited amount of data available on electronic device footprint we were not able to capture every single component. We were able to account for 8 out of 11 components from our bill of materials, though, which we feel captures the concept sufficiently for the sake of demonstration. We were unable to find data for the connectors and voltage regulator. However, in addition to the bill of materials, we capture the carbon footprint for the ML sensor's model training, transport, and three-year use.

The total carbon footprint, including embodied and operational footprint, of our ML Sensor is approximately **2.82 kg CO2-eq**. The chart below shows how the footprint is broken down. The majority of the footprint can be attributed to the power supply and camera sensor.



We note that we do not claim that this is 100% accurate but rather a representative approximation of the sensor's environmental impact and what other future datasheet should aim to include.

Acronyms

Acronym	Description	
GDPR	General Data Protection Regulation	
GPIO	General Purpose Input Output	
ML	Machine Learning	
I2C	Inter-Integrated Circuit	
ID	Identifier	
IoU	Intersection Over Union	
LED	Light-Emitting Diode	
MCU	Microcontroller Unit	
MEMS	Microelectromechanical System	
MST	Monk Skin Tone Scale	
PCA	Principal Component Analysis	
RGB	Red Blue Green	

Glossary

Lux Photometric unit of luminance (at 550 nm, 1 lux = 1 lumen/ m^2 = 1/683 W/ m^2)

Sensitivity A measure of pixel performance that characterizes the rise of the photodiode or sense node signal in Volts

upon illumination with light. Units are typically $V/(W/m^2)$ /sec and are dependent on the incident light wavelength. Sensitivity measurements are often taken with 550 nm incident light. At this wavelength, 683 lux is equal to 1 W/m^2 ; the units of sensitivity are quoted in V/lux/sec. Note that responsivity and sensitivity are used interchangeably in image sensor characterization literature so it is best to check the

units.

IoU Intersection Over Union (IoU) is a metric used to evaluate the accuracy of an object detector on a specific

dataset. It measures the overlap between the predicted bounding box (from the detector) and the ground truth bounding box. Values range between 0 and 1, where a higher value indicates better prediction accuracy. A value of 0 indicates no overlap, while a value of 1 indicates perfect overlap (the predicted box

matches the ground truth exactly).

Inference The process of applying a trained machine learning model to unseen data for making predictions or

classifications. In the context of person detection, it involves analyzing images or video frames to

determine if a person is present.

Accuracy A performance metric that measures the overall correctness of a person detection system, indicating the

percentage of correctly identified persons in the total number of instances.

Principal Component Analysis A statistical procedure that transforms a set of observations of possibly correlated variables into a set of values of linearly uncorrelated variables called principal components. These components are orthogonal to

each other and capture the variance in the data in decreasing order.

Monk Skin Tone Scale A 10-shade system, developed by Google, designed to provide a more inclusive representation of diverse skin tones in image-based technologies to address the challenges of representation in image-based

technologies, especially for people of color.

Training Set Labeled examples or samples used to teach a machine learning model to recognize and classify objects

accurately. In the case of person detection, it comprises images or videos with annotated information about

the presence or absence of people.

Test Set A subset of the dataset that is strictly used to evaluate the performance of a model after it has been trained.

The test set provides an unbiased evaluation of a model's generalization to new, unseen data. It should

never be used during training or hyperparameter tuning.

Validation Set A subset of the dataset, separate from the training set, used to evaluate a model during training. It provides

an intermittent check on the model's performance, allowing for hyperparameter tuning and model selection. By evaluating model performance on a validation set, one can detect issues like overfitting (where the model performs exceptionally well on the training set but poorly on new, unseen data). Once

the model is optimized using the validation set, its final performance is then assessed on the test set.

Person Detection The process of identifying the presence and location of a person within an image or video stream.

Sensor A device that detects and measures physical or environmental properties, such as the presence of a person,

and converts them into electrical signals.